

```
In [94]: # Core Libraries
import pandas as pd
import numpy as np

# Visualization Libraries
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

# Statistical and Machine Learning Libraries
from scipy import stats
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score

# Time-Series Analysis Libraries (if needed)
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.arima.model import ARIMA

# Geographic Data Libraries (if needed)
import geopandas as gpd
import folium
from shapely.geometry import Point, Polygon

# Data Cleaning and Transformation
from datetime import datetime

In [93]: !pip install pandas numpy matplotlib seaborn scikit-learn statsmodels plotly geopandas
!pip install folium

Requirement already satisfied: pandas in c:\users\adesh\anaconda3\lib\site-packages (2.1.4)
Requirement already satisfied: numpy in c:\users\adesh\anaconda3\lib\site-packages (1.26.4)
Requirement already satisfied: matplotlib in c:\users\adesh\anaconda3\lib\site-packages (3.8.0)
Requirement already satisfied: seaborn in c:\users\adesh\anaconda3\lib\site-packages (0.12.2)
Requirement already satisfied: scikit-learn in c:\users\adesh\anaconda3\lib\site-packages (1.2.2)
Requirement already satisfied: statsmodels in c:\users\adesh\anaconda3\lib\site-packages (0.14.0)
Requirement already satisfied: plotly in c:\users\adesh\anaconda3\lib\site-packages (5.9.0)
Requirement already satisfied: geopandas in c:\users\adesh\anaconda3\lib\site-packages (1.0.1)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\adesh\anaconda3\lib\site-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in c:\users\adesh\anaconda3\lib\site-packages (from pandas) (2023.3.post1)
Requirement already satisfied: tzdata>=2022.1 in c:\users\adesh\anaconda3\lib\site-packages (from pandas) (2023.3)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\adesh\anaconda3\lib\site-packages (from matplotlib) (1.2.0)
Requirement already satisfied: cycler>=0.10 in c:\users\adesh\anaconda3\lib\site-packages (from matplotlib) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\adesh\anaconda3\lib\site-packages (from matplotlib) (4.25.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\adesh\anaconda3\lib\site-packages (from matplotlib) (1.4.4)
Requirement already satisfied: packaging>=20.0 in c:\users\adesh\anaconda3\lib\site-packages (from matplotlib) (23.1)
Requirement already satisfied: pillow>=6.2.0 in c:\users\adesh\anaconda3\lib\site-packages (from matplotlib) (10.2.0)
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\adesh\anaconda3\lib\site-packages (from matplotlib) (3.0.9)
Requirement already satisfied: scipy>=1.2 in c:\users\adesh\anaconda3\lib\site-packages (from scikit-learn) (1.11.4)
Requirement already satisfied: joblib>=1.1.1 in c:\users\adesh\anaconda3\lib\site-packages (from scikit-learn) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\adesh\anaconda3\lib\site-packages (from scikit-learn) (2.2.0)
Requirement already satisfied: patsy>=0.5.2 in c:\users\adesh\anaconda3\lib\site-packages (from statsmodels) (0.5.3)
Requirement already satisfied: tenacity>=6.2.0 in c:\users\adesh\anaconda3\lib\site-packages (from plotly) (8.2.2)
Requirement already satisfied: pyogrio>=0.7.2 in c:\users\adesh\anaconda3\lib\site-packages (from geopandas) (0.10.0)
Requirement already satisfied: pyproj>=3.3.0 in c:\users\adesh\anaconda3\lib\site-packages (from geopandas) (3.7.0)
Requirement already satisfied: shapely>=2.0.0 in c:\users\adesh\anaconda3\lib\site-packages (from geopandas) (2.0.6)
Requirement already satisfied: six in c:\users\adesh\anaconda3\lib\site-packages (from patsy>=0.5.2->statsmodels) (1.16.0)
Requirement already satisfied: certifi in c:\users\adesh\anaconda3\lib\site-packages (from pyogrio>=0.7.2->geopandas) (2024.2.2)
Collecting folium
  Downloading folium-0.19.2-py2.py3-none-any.whl.metadata (3.8 kB)
Collecting branca>=0.6.0 (from folium)
  Downloading branca-0.8.1-py3-none-any.whl.metadata (1.5 kB)
Requirement already satisfied: Jinja2>=2.9 in c:\users\adesh\anaconda3\lib\site-packages (from folium) (3.1.3)
Requirement already satisfied: MarkupSafe in c:\users\adesh\anaconda3\lib\site-packages (from Jinja2>=2.9->folium) (2.1.3)
Requirement already satisfied: requests in c:\users\adesh\anaconda3\lib\site-packages (from requests>=follum) (2.0.4)
Requirement already satisfied: idna<4,>=2.5 in c:\users\adesh\anaconda3\lib\site-packages (from requests>=follum) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\adesh\anaconda3\lib\site-packages (from requests>=follum) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in c:\users\adesh\anaconda3\lib\site-packages (from requests>=follum) (2024.2.2)
Downloading folium-0.19.2-py2.py3-none-any.whl (110 kB)
----- 0.0/110.1 kB 7 eta -:--:--
----- 30.7/110.1 kB 660.6 kB/s eta 0:00:01
----- 102.4/110.1 kB 2.0 MB/s eta 0:00:00
----- 110.1/110.1 kB 1.3 MB/s eta 0:00:00
Downloading branca-0.8.1-py3-none-any.whl (26 kB)
Installing collected packages: branca, folium
Successfully installed branca-0.8.1 folium-0.19.2

In [11]: #df = pd.read_csv(r"C:\Users\Adesh\Downloads\Air_Quality.csv")

In [6]: print(df.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18025 entries, 0 to 18024
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype
---  --
 0   Unique ID   18025 non-null  int64
 1   Indicator ID 18025 non-null  int64
 2   Name        18025 non-null  object
 3   Measure     18025 non-null  object
 4   Measure Info 18025 non-null  object
 5   Geo Type Name 18025 non-null  object
 6   Geo Join ID  18016 non-null  float64
 7   Geo Place Name 18016 non-null  object
 8   Time Period  18025 non-null  object
 9   Start_Date  18025 non-null  object
10   Data Value   18025 non-null  float64
11   Message      0 non-null      object
dtypes: float64(3), int64(2), object(7)
memory usage: 1.7+ MB
None

In [7]: print(df.isnull().sum())

Unique ID      0
Indicator ID    0
Name            0
Measure         0
Measure Info    0
Geo Type Name   0
Geo Join ID     9
Geo Place Name  9
Time Period     0
Start_Date      0
Data Value      0
Message         18025
dtype: int64

In [14]: # Remove rows where 'Geo Join ID' or 'Geo Place Name' is null
df_cleaned = df.dropna(subset=['Geo Join ID', 'Geo Place Name'])
df_cleaned = df.drop(columns=['Message'])

# Verify the updated dataset
print(f"Rows before cleaning: {len(df)}")
print(f"Rows after cleaning: {len(df_cleaned)}")
print(f"Columns before cleaning: {df.columns.tolist()}")
print(f"Columns after cleaning: {df_cleaned.columns.tolist()}")

# Save the cleaned dataset to a new file (optional)
df_cleaned.to_csv(r"C:\Users\Adesh\Downloads\Air_Quality_Cleaned.csv", index=False)

Rows before cleaning: 18016
Rows after cleaning: 18016
Columns before cleaning: ['Unique ID', 'Indicator ID', 'Name', 'Measure', 'Measure Info', 'Geo Type Name', 'Geo Join ID', 'Geo Place Name', 'Time Period', 'Start_Date', 'Data Value', 'Message']
Columns after cleaning: ['Unique ID', 'Indicator ID', 'Name', 'Measure', 'Measure Info', 'Geo Type Name', 'Geo Join ID', 'Geo Place Name', 'Time Period', 'Start_Date', 'Data Value']

In [15]: df = pd.read_csv(r"C:\Users\Adesh\Downloads\Air_Quality_Cleaned.csv")

In [17]: #Standardize Date Format
df['Start_Date'] = pd.to_datetime(df['Start_Date'], format='%m/%d/%Y')

In [22]: # Standardize Text Column Strip whitespace and convert to lowercase
df['Geo Place Name'] = df['Geo Place Name'].str.strip().str.lower()

In [59]: df = df[df['Time Period'].str.contains("Annual Average")]

In [60]: df['Time Period'] = df['Time Period'].str.replace("Annual Average", "").str.strip()

In [61]: print(df.describe())

count      3948.000000    3948.000000    3.948000e+03
mean      418589.597518      370.000000    7.622145e+05
min       167509.000000      365.000000    1.000000e+00
25%       176565.750000      365.000000    2.030000e+02
50%       373808.500000      370.000000    3.030000e+02
75%       645204.250000      375.000000    4.040000e+02
max       826379.000000      375.000000    1.051061e+08
std       236987.987911      5.006633    8.820152e+06

              Start_Date  Data Value  Month
count      2015-06-19 10:17:08.571428352    14.340502    6.500000
mean      2008-12-01 00:00:00      5.000000    1.000000
25%      2011-12-01 00:00:00      8.400000    1.000000
50%      2015-07-02 00:00:00     12.100000    6.500000
75%      2019-01-01 00:00:00     19.700000    12.000000
max      2022-01-01 00:00:00     46.800000    12.000000
std              NaN      7.004176    5.500697

In [63]: #Unique Values in Categorical Column
print(df['Geo Type Name'].unique())
print(df['Name'].unique())

['URB42' 'Borough' 'URB34' 'CD' 'Citywide']
['Fine particles (PM 2.5)' 'Nitrogen dioxide (NO2)']

In [64]: df['Data Value'].hist(bins=50, range=(0,50) )
plt.title("Distribution of Air Quality Values")
plt.xlabel("Data Value")
plt.ylabel("Frequency")
plt.show()

In [65]: # Create a line plot for seasonal trends
seasonal_trends = df.groupby('Month')['Data Value'].mean()
seasonal_trends.plot(kind='line', marker='o', linestyle='-', color='green', title="Seasonal Air Quality Trends")

plt.xlabel("Month")
plt.ylabel("Average Data Value")
plt.xticks(ticks=seasonal_trends.index) # Ensures proper month labeling
plt.grid(True) # Add grid lines for better readability
plt.show()

In [88]: # Correlation heatmap
le = LabelEncoder()
df['Name_encoded'] = le.fit_transform(df['Name'])

# Re-run heatmap with numeric columns
df = df.select_dtypes(include=['float64', 'int64'])
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()

In [90]: # Create column 'Air Quality category' (low, moderate, high)
def categorize_quality(value):
    if value < 10:
        return 'Low'
    elif 10 <= value < 20:
        return 'Moderate'
    else:
        return 'High'

df['Quality Category'] = df['Data Value'].apply(categorize_quality)
print(df['Quality Category'].value_counts())

Quality Category
Low      1564
Moderate 1462
High     822
Name: count, dtype: int64

In [67]: df.to_csv(r"C:\Users\Adesh\Downloads\Air_Quality_Cleaned.csv", index=False)

In [78]: # Advanced Predictive Modeling
# Prepare the data
X = df[['Geo Join ID', 'Indicator ID', 'Time Period']] # Features
y = df['Data Value']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

encoder = LabelEncoder()
df['Geo Place Name Encoded'] = encoder.fit_transform(df['Geo Place Name'])

scaler = StandardScaler()
df_scaled = scaler.fit_transform(df[['Geo Join ID', 'Indicator ID', 'Geo Place Name Encoded']])

# Define input features (X) and target variable (y)
X = df[['Geo Join ID', 'Indicator ID', 'Geo Place Name Encoded']]
y = df['Data Value']

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

In [81]: # Initialize the model
model = RandomForestRegressor(n_estimators=100, random_state=42)

# Train the model
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-Squared: {r2}")

Mean Squared Error: 6.526543563388008
R-Squared: 0.866858944523164

In [ ]:

In [46]: #load data into SQL server
import sqlalchemy as sa
engine = sa.create_engine('mysql://AADESH@SQLEXPRESS/master:driver=ODBCDRIVER17+FOR+SQL+SERVER')
conn=engine.connect()

In [47]: # provide connection to SQL server
df.to_sql('df_NYC_Air_Quality', conn=conn, index=False, if_exists='replace')

Out[47]: 145
```

